

1. Image Classification Pipeline

- An image is a big grid of numbers between $[0, 255]$

$800 \times 600 \times 3$

↳ 3 channels of RGB

$32 \times 32 \times 3 = 3072$ numbers in an image

- First classifier: Nearest Neighbor

- Distance Metric to compare imgs

L_1 distance: $d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$ # Manhattan

- K-Nearest Neighbors

L_2 distance: $d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$ # Euclidean

- Decide on the best value of k and distance

- Split data into train, validation and test


Choose hyperparameters on validation & evaluate

on test

only once at the end


Cross-validation

Linear Classification



Array of $32 \times 32 \times 3$ numbers
(3072 numbers)

e.g. 4 pixels, ~~factor~~ only R channel
3 classes



$$f(x, W) = Wx + b \rightarrow 10 \times 1$$

10×1 10×3072 3072×1

$$f(x, W) = Wx + b$$

3×4 4×1 3×1

• Loss function:

tells how good current classifier is

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

actual label (x_i, y_i) : x_i is the img

$S = f(x_i, W)$ & y_i is the integer label

scores vector

QVM Loss

$$L_i = \sum_{j \neq y_i} \begin{cases} 0 & \text{if } S_{y_i} \geq S_j + 1 \\ S_j - S_{y_i} + 1 & \text{otherwise} \end{cases}$$

$$L = \frac{1}{N} \sum_{i=1}^N L_i$$

$$= \sum_{j \neq y_i} \max(0, S_j - S_{y_i} + 1)$$

正确分数比当前分数高，则增加正确分数

Regularization

$$L(W) = \frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i) + \lambda R(W)$$

prevent model from
doing too well on training data
prefers simpler models

L2 Regularization: $R(W) = \sum_k \sum_l W_{k,l}^2$

L1 Regularization: $R(W) = \sum_k \sum_l |W_{k,l}|$

Softmax Classifier (Multinomial Logistic Regression)

$$P(Y=k | X=x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

$$L_i = -\log P(Y=y_i | X=x_i)$$

最大似然估计

1. 交叉熵 (Kullback-Leibler divergence)

$$D_{KL}(P || Q) = \sum_y P(y) \log \frac{P(y)}{Q(y)}$$

2. 交叉熵 (Cross Entropy) ★

$$H(P, Q) = H(P) + D_{KL}(P || Q)$$

1. e.g.

0.13	compare	1.00
0.87	→ ←	0
0.00	K-L	0
	P Q	

Correct

2. e.g.

Same as 1

Summary :

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

Softmax

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

SIM

$$L = \frac{1}{N} \sum_{i=1}^N L_i + R(W)$$

Full Loss

$$R(W) = \sum_k W_k^2$$

Optimization | find the best W for best loss

1. Random search & update W & loss

2. Gradient descent

$$\nabla_W L(W) = \frac{1}{N} \sum_{i=1}^N \nabla_W L_i(x_i, y_i, W) + \lambda \nabla_W R(W)$$

$$W += -\text{step-size} * \nabla_W$$

Image Features

Build $\{$ Extract random patches

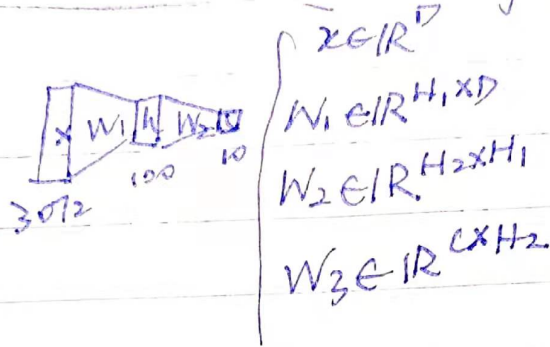
codebook \cdot Cluster patches to form codebook of visual words

Encode

Images


Neural Networks & Backpropagation

2-layer Neural Network: $f = W_2 \max(0, W_1 x)$
 3-layer: $f = W_3 \max(0, W_2 \max(0, W_1 x))$

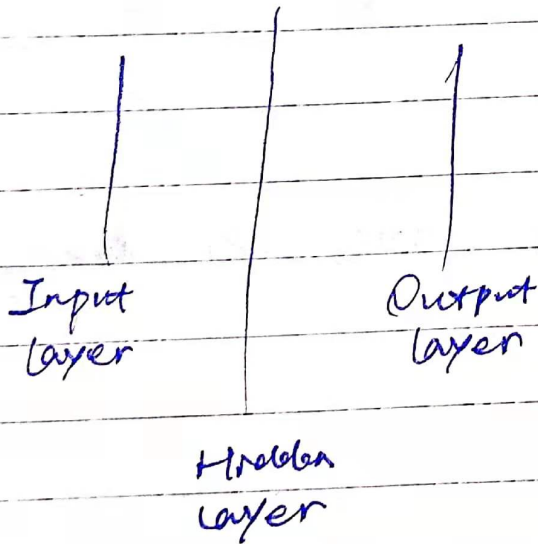


$\max(0, z) = \text{activation function}$

ReLU:



Other activation functions: Sigmoid, tanh, ReLU (default).



2-layer NN

Forward pass of a NN

$f = \# \text{sigmoid}$

$x = \text{np.random.randn}(3, 1)$ # nodes input

$h_1 = f(\text{np.dot}(W_1, x) + b_1)$

$h_2 = f(\text{np.dot}(W_2, h_1) + b_2)$

$\text{out} = \text{np.dot}(W_3, h_2) + b_3$

$x: 3 \times 1$

$W_1: 4 \times 3$

$W_2: 4 \times 4$

$W_3: 1 \times 4$

$$L = \frac{1}{N} \sum_{i=1}^N L_i + \lambda R(W_1) + \lambda R(W_2)$$

★ Back propagation to compute $\frac{\partial L}{\partial W_1}$, $\frac{\partial L}{\partial W_2}$

e.g. $f = (x+y)z$ $\Rightarrow \frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial y}$
 $q = x+y$

↑ ↓
 Upstream gradient Local gradient

Forward pass: compute output

Backward pass: compute grads

Vector Derivatives:

$x \in \mathbb{R}^N$, $y \in \mathbb{R}^m$

$\frac{\partial y}{\partial x} \in \mathbb{R}^{N \times m}$

$$\left(\frac{\partial y}{\partial x} \right)_{n,m} = \frac{\partial y_m}{\partial x_n}$$

$$\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

$$\begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \end{bmatrix}$$

$N=3$

$m=4$

$\mathbb{R}^{3 \times 4}$

$$\begin{bmatrix} - & - & - & - \\ - & - & - & - \\ - & - & - & - \end{bmatrix}$$

off-diagonal entries always 0

Jacobian * upstream gradient

$$\frac{dL}{dx} = \frac{dy}{dx} \cdot \frac{dL}{dy}$$

$$\left(\frac{\partial L}{\partial x} \right)_i = \begin{cases} \left(\frac{\partial L}{\partial y} \right)_i & \text{if } x_i = y_i \\ 0 & \text{e.w.} \end{cases}$$

Jacobian e.g. $f(x, y, z) = (xy + 2yz, 2xy^2z)$

$$\begin{bmatrix} y & x+2z & 2y \\ 2yz & 4xy/z & 2xy^2 \end{bmatrix} \begin{matrix} \rightarrow xy + 2yz \text{ w.r.t } x, y, z \\ \rightarrow 2xy^2z \text{ w.r.t } x, y, z \end{matrix}$$

No Explicit
 Jacobian ~~is~~: Memory not enough

(2)

$$X = \begin{bmatrix} x_{11} & x_{12} \\ x_{21} & x_{22} \end{bmatrix}$$

$$W = \begin{bmatrix} w_{11} & - & - & w_{13} \\ - & - & - & w_{23} \end{bmatrix}$$

$$\frac{\partial L}{\partial x} : \begin{bmatrix} - & - \\ - & - \end{bmatrix} \quad Y = XW = \begin{bmatrix} - & - \\ - & - \end{bmatrix}_{2 \times 3} \quad \frac{\partial L}{\partial Y} : \begin{bmatrix} - & - \\ - & - \end{bmatrix}_{2 \times 3}$$

$$\frac{\partial L}{\partial x_{11}} = \frac{\partial L}{\partial Y} \cdot \frac{\partial Y}{\partial x_{11}}$$

★ dot product to get a scalar

same shape as Y

$$\begin{cases} \frac{\partial L}{\partial x} = \frac{\partial L}{\partial Y} W^T \\ \frac{\partial L}{\partial W} = X^T \frac{\partial L}{\partial Y} \end{cases}$$